

# Design Decision

This project follows production-grade architectural choices to ensure scalability, performance, security, and maintainability.

## *Database: PostgreSQL with Timescale DB*

### Why Timescale DB?

Timescale DB is purpose-built for time-series workloads, making it ideal for high-volume API traffic logs.

- Time-series optimization  
Automatic partitioning using hyper tables for efficient time-based queries.
- SQL compatibility  
Fully compatible with PostgreSQL — no new query language required.
- Compression  
Built-in compression reduces storage usage by up to 90%.
- High performance  
Optimized for both high write throughput (log ingestion) and fast analytics queries.
- Production maturity  
Battle-tested with strong community support and documentation.

### Schema Design

- Hypertable partitioned on the timestamp column
- Composite indexes:
  - (timestamp, status\_code)
  - (timestamp, service\_name)
- Additional indexes for efficient filtering of:
  - 2xx, 4xx, and 5xx HTTP status code ranges

## *Backend: Node.js (Express) with TypeScript*

### Why This Stack?

The backend is designed for high concurrency, type safety, and data integrity.

- High throughput  
Node.js event-driven architecture handles concurrent API requests efficiently.
- Type safety  
TypeScript reduces runtime errors and improves long-term maintainability.
- Input validation  
Zod schemas enforce strict validation at API boundaries.
- Connection pooling  
pg-pool efficiently manages PostgreSQL connections under load.

## *Containerization: Docker Compose*

### Architecture Decisions

Docker Compose is used to orchestrate all services in a consistent, reproducible environment.

- Multi-stage builds  
Produces optimized and lightweight production images.
- Named volumes  
Ensures database data persists across container restarts.
- Health checks  
Services wait for dependencies before starting.
- Network isolation  
Internal service communication is isolated; only Nginx is exposed externally.

## *Security Implementation*

Feature	Description
API Authentication	API key validation via x-api-key header (bcrypt-hashed)
Rate Limiting	express-rate-limit + Nginx rate limiting
Input Validation	Zod schemas for all request payloads
SQL Injection Protection	Parameterized SQL queries
CORS	Restricted to configured origins
Request Size Limits	1MB (backend), 2MB (Nginx)
Security Headers	X-Frame-Options, X-Content-Type-Options, X-XSS-Protection

## *Performance Optimizations*

Optimization	Details
Database Indexing	Composite indexes on frequently queried columns
Connection Pooling	Maximum 20 connections, 30s idle timeout
Batch Ingestion	Single INSERT for up to 1000 log entries
Query Optimization	Time-based aggregations using time_bucket()
Static Asset Caching	Enabled via Nginx
Gzip Compression	Enabled for all text-based responses